

Brief

Write the architecture of a SimpleUNet following the diagram shown during the lecture/in GitHub:

Encoder: Define a convolutional block (conv_block) that will be repeated in each layer of the encoder. Define as well a maxpooling block (maxpool_block) that will also be repeated when necessary during the encoding (downsampling path). Hint: use torch.nn.Conv2d and remember you need to add by hand the activation function (use ReLU).

Middle: define the middle or bottleneck part of the architecture.

Decoder: Define the transposed convolutional blocks (transposed_block). Hint: use nn.ConvTranspose2d.

Final layer: Define the final convolutional layer.

Forward: Remember you need to define the forward pass of your class, with how each block will be called.

In order to aid the optimisation process, include batch normalisation (nn.BatchNorm2d) and 50% dropout (nn.Dropout2d(0.5)). Where do think they should be placed?

Note: There are different ways to build this architecture, here we will use PyTorch as we will test that we can load the pre-trained model provided.

```
import torch
import torch.nn as nn
from torch.utils.data import TensorDataset, DataLoader
import matplotlib.pyplot as plt
!pip install torchmetrics
import torchmetrics
from torchmetrics.classification import BinaryAccuracy
import numpy as np
from torchsummary import summary
import cv2 # optional for the bonus questions
```

```
Collecting torchmetrics
  Downloading torchmetrics-1.3.1-py3-none-any.whl (840 kB)
    840.4/840.4 kB 5.5 MB/s eta 0:00:00
Requirement already satisfied: numpy>1.20.0 in /usr/local/lib/python3.10/dist-packages (from torchmetrics) (1.25.2)
Requirement already satisfied: packaging>17.1 in /usr/local/lib/python3.10/dist-packages (from torchmetrics) (23.2)
Requirement already satisfied: torch>=1.10.0 in /usr/local/lib/python3.10/dist-packages (from torchmetrics) (2.1.0+cu121)
Collecting lightning-utilities>=0.8.0 (from torchmetrics)
  Downloading lightning_utilities-0.10.1-py3-none-any.whl (24 kB)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from lightning-utilities>=0.8.0->torchmetrics) (58.1.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from lightning-utilities>=0.8.0->torchmetrics) (4.5.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (3.12.2)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (1.11.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (3.1)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (3.1.2)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (2023.12.0)
Requirement already satisfied: triton==2.1.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->torchmetrics) (2.1.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.10.0->torchmetrics) (2.1.2)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch>=1.10.0->torchmetrics) (1.3.0)
Installing collected packages: lightning-utilities, torchmetrics
Successfully installed lightning-utilities-0.10.1 torchmetrics-1.3.1
```

Fill the empty spaces to define the architecture of the UNet as the image provided in the problem sheet (call it SimpleUNet)

```
class SimpleUNet(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.conv1 = self.conv_block(in_channels, 16, 3, 1, 1)
        self.maxpool1 = self.maxpool_block(2, 2, 0)
        self.conv2 = self.conv_block(16, 32, 3, 1, 1)
        self.maxpool2 = self.maxpool_block(2, 2, 0)
        self.conv3 = self.conv_block(32, 64, 3, 1, 1)
        self.maxpool3 = self.maxpool_block(2, 2, 0)

        self.middle = self.conv_block(64, 128, 3, 1, 1)

        self.upsample3 = self.transposed_block(128, 64, 3, 2, 1, 1)
        self.upconv3 = self.conv_block(128, 64, 3, 1, 1)
        self.upsample2 = self.transposed_block(64, 32, 3, 2, 1, 1)
        self.upconv2 = self.conv_block(64, 32, 3, 1, 1)

        self.upsample1 = self.transposed_block(32, 16, 3, 2, 1, 1)
        self.upconv1 = self.conv_block(32, 16, 3, 1, 1)

        self.final = self.final_layer(16, 1, 1, 1, 0)

    def conv_block(self, in_channels, out_channels, kernel_size, stride, padding):
        convolution = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=kernel_size, stride=stride, padding=padding),
```

```

        nn.BatchNorm2d(out_channels),
        nn.ReLU(inplace=True),
        nn.Conv2d(out_channels, out_channels, kernel_size=kernel_size, stride=stride, padding=padding),
        nn.BatchNorm2d(out_channels),
        nn.ReLU(inplace=True))

    return convolution

def maxpool_block(self, kernel_size, stride, padding):
    maxpool = nn.Sequential(nn.MaxPool2d(kernel_size=kernel_size, stride=stride, padding=padding),
                            nn.Dropout2d(0.5))
    return maxpool

def transposed_block(self, in_channels, out_channels, kernel_size, stride, padding, output_padding):
    transposed = torch.nn.ConvTranspose2d(in_channels, out_channels, kernel_size=kernel_size, stride=stride,
                                           padding=padding, output_padding=output_padding)

    return transposed

def final_layer(self, in_channels, out_channels, kernel_size, stride, padding):
    final = nn.Conv2d(in_channels, out_channels, kernel_size=kernel_size, stride=stride, padding=padding)
    return final

def forward(self, x):
    # downsampling part
    conv1 = self.conv1(x)
    maxpool1 = self.maxpool1(conv1)
    conv2 = self.conv2(maxpool1)
    maxpool2 = self.maxpool2(conv2)
    conv3 = self.conv3(maxpool2)
    maxpool3 = self.maxpool3(conv3)

    # middle part
    middle = self.middle(maxpool3)

    # upsampling part
    upsample3 = self.upsample3(middle)
    upconv3 = self.upconv3(torch.cat([upsample3, conv3], 1))
    upsample2 = self.upsample2(upconv3)
    upconv2 = self.upconv2(torch.cat([upsample2, conv2], 1))
    upsample1 = self.upsample1(upconv2)
    upconv1 = self.upconv1(torch.cat([upsample1, conv1], 1))

    final_layer = self.final(upconv1)

    return final_layer

```

Now load the pre-trained model and test if the implementation above is correct by printing the summary of the model.

```

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
model = SimpleUNet(in_channels=1, out_channels=1)
model.load_state_dict(torch.load('SimpleUNet_v3.pt', map_location=torch.device(device)))
summary(model, input_size=(1, 512, 512))

```

Layer (type)	Output Shape	Param #

Conv2d-1	[-1, 16, 512, 512]	160
BatchNorm2d-2	[-1, 16, 512, 512]	32
ReLU-3	[-1, 16, 512, 512]	0
Conv2d-4	[-1, 16, 512, 512]	2,320
BatchNorm2d-5	[-1, 16, 512, 512]	32
ReLU-6	[-1, 16, 512, 512]	0
MaxPool2d-7	[-1, 16, 256, 256]	0
Dropout2d-8	[-1, 16, 256, 256]	0
Conv2d-9	[-1, 32, 256, 256]	4,640
BatchNorm2d-10	[-1, 32, 256, 256]	64
ReLU-11	[-1, 32, 256, 256]	0
Conv2d-12	[-1, 32, 256, 256]	9,248
BatchNorm2d-13	[-1, 32, 256, 256]	64
ReLU-14	[-1, 32, 256, 256]	0
MaxPool2d-15	[-1, 32, 128, 128]	0
Dropout2d-16	[-1, 32, 128, 128]	0
Conv2d-17	[-1, 64, 128, 128]	18,496
BatchNorm2d-18	[-1, 64, 128, 128]	128
ReLU-19	[-1, 64, 128, 128]	0
Conv2d-20	[-1, 64, 128, 128]	36,928
BatchNorm2d-21	[-1, 64, 128, 128]	128
ReLU-22	[-1, 64, 128, 128]	0
MaxPool2d-23	[-1, 64, 64, 64]	0
Dropout2d-24	[-1, 64, 64, 64]	0
Conv2d-25	[-1, 128, 64, 64]	73,856
BatchNorm2d-26	[-1, 128, 64, 64]	256
ReLU-27	[-1, 128, 64, 64]	0

Conv2d-28	[-1, 128, 64, 64]	147,584
BatchNorm2d-29	[-1, 128, 64, 64]	256
ReLU-30	[-1, 128, 64, 64]	0
ConvTranspose2d-31	[-1, 64, 128, 128]	73,792
Conv2d-32	[-1, 64, 128, 128]	73,792
BatchNorm2d-33	[-1, 64, 128, 128]	128
ReLU-34	[-1, 64, 128, 128]	0
Conv2d-35	[-1, 64, 128, 128]	36,928
BatchNorm2d-36	[-1, 64, 128, 128]	128
ReLU-37	[-1, 64, 128, 128]	0
ConvTranspose2d-38	[-1, 32, 256, 256]	18,464
Conv2d-39	[-1, 32, 256, 256]	18,464
BatchNorm2d-40	[-1, 32, 256, 256]	64
ReLU-41	[-1, 32, 256, 256]	0
Conv2d-42	[-1, 32, 256, 256]	9,248
BatchNorm2d-43	[-1, 32, 256, 256]	64
ReLU-44	[-1, 32, 256, 256]	0
ConvTranspose2d-45	[-1, 16, 512, 512]	4,624
Conv2d-46	[-1, 16, 512, 512]	4,624
BatchNorm2d-47	[-1, 16, 512, 512]	32
ReLU-48	[-1, 16, 512, 512]	0
Conv2d-49	[-1, 16, 512, 512]	2,320
BatchNorm2d-50	[-1, 16, 512, 512]	32
ReLU-51	[-1, 16, 512, 512]	0
Conv2d-52	[-1, 1, 512, 512]	17

=====

Total params: 536,913

Test applying the model to a few 2D slices of the test case

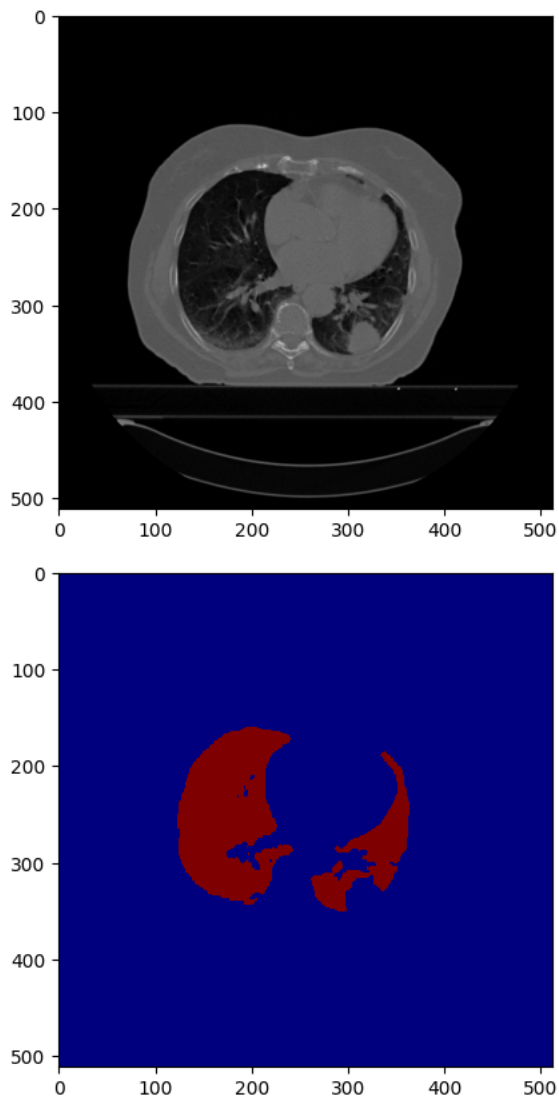
```
# Load the images to test
batch_size = 1
image_file = 'Case_010.npz'

image_array = (np.load(image_file))['images']
image_tensor = torch.Tensor(image_array)

# Model to evaluation mode
model.eval()
metric = BinaryAccuracy()
threshold = 0.5

# Choose and plot a 2D slice
slice_id = 42
curr_image = image_tensor[slice_id, :, :]
plt.imshow(curr_image, cmap='gray')
plt.show()




# We need to convert our input to 4D (including batch size) to compute predictions
curr_image = torch.unsqueeze(curr_image, 0)
curr_image = torch.unsqueeze(curr_image, 0)
curr_pred = model(curr_image)
curr_pred = torch.sigmoid(curr_pred)
curr_pred = (curr_pred > threshold)
curr_pred = curr_pred[0,0,:,:]
plt.imshow(curr_pred, cmap='jet', interpolation='none')
plt.show()
```



BONUS: Apply the Sobel algorithms using OpenCV

```
# Get again the slice as a numpy array
curr_slice = image_array[slice_id, :, :]

# Sobel Edge Detection
sobelx = cv2.Sobel(src=curr_slice, ddepth=cv2.CV_64F, dx=1, dy=0, ksize=5) # Sobel Edge Detection on the X axis
sobely = cv2.Sobel(src=curr_slice, ddepth=cv2.CV_64F, dx=0, dy=1, ksize=5) # Sobel Edge Detection on the Y axis
sobelxy = cv2.Sobel(src=curr_slice, ddepth=cv2.CV_64F, dx=1, dy=1, ksize=5) # Combined X and Y Sobel Edge Detection
# Display Sobel Edge Detection Images
plt.imshow(sobelx, cmap='gray')
plt.show()
plt.imshow(sobely, cmap='gray')
plt.show()
plt.imshow(sobelxy, cmap='gray')
plt.show()
```



Compare to Thresholding results: binarisation and Otsu



```
image_file = 'Case_010.npz'
image_array = (np.load(image_file))['images']
slice_id = 32
curr_slice = image_array[slice_id, :, :]

# Binarisation:
# we will use the knowledge that air is at -1000 H.U.
# in this case we haven't converted to H.U. yet
curr_slice_threshold = curr_slice > 400
plt.imshow(curr_slice_threshold, cmap='gray')
plt.show()

# Otsu, using OpenCV, and an uint8 version of the image
curr_slice_int = curr_slice.astype("uint8")
ret,image_otsu = cv2.threshold(curr_slice_int,0,255,cv2.THRESH_BINARY+cv2.THRESH_OTSU)
plt.imshow(image_otsu, cmap='gray')
plt.show()
```

